

Project : Health care

Build a model to accurately predict whether the patients in the dataset have diabetes or not?

Context:

This dataset is originally from the National Institute of Diabetes and Digestive and Kidney Diseases. The objective of the dataset is to diagnostically predict whether or not a patient has diabetes, based on certain diagnostic measurements included in the dataset. Several constraints were placed on the selection of these instances from a larger database. In particular, all patients here are females at least 21 years old of Pima Indian heritage.

Problem Statement:

Build a model to accurately predict whether the patients in the dataset have diabetes or not?

Dataset Description:

The datasets consists of several medical predictor variables and one target variable, Outcome. Predictor variables includes the number of pregnancies the patient has had, their BMI, insulin level, age, and so on.

Pregnancies: Number of times pregnant

Glucose: Plasma glucose concentration a 2 hours in an oral glucose tolerance test

BloodPressure: Diastolic blood pressure (mm Hg)

SkinThickness: Triceps skin fold thickness (mm)

Insulin: 2-Hour serum insulin (mu U/ml)

BMI: Body mass index (weight in kg/(height in m)^2)

DiabetesPedigreeFunction: Diabetes pedigree function

Age: Age (years)

Outcome: Class variable (0 or 1) 268 of 768 are 1, the others are 0

```
In [1]: # Import the required Libraries
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split , cross_val_score
from sklearn.model_selection import GridSearchCV
```

```

from sklearn.linear_model import LogisticRegression
from sklearn.neighbors import KNeighborsClassifier
from sklearn.svm import SVC
from sklearn.tree import DecisionTreeClassifier
from sklearn import metrics
from sklearn.metrics import roc_curve, auc
from sklearn.ensemble import RandomForestClassifier, BaggingClassifier, AdaBoostClassifier

import warnings
warnings.filterwarnings('ignore')

```

```

In [2]: ## Read the dataset
healthcare=pd.read_csv('health care diabetes.csv')

```

```

In [3]: healthcare.head()

```

```

Out[3]:

```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction	Outcome
0	6	148	72	35	0	33.6	0.627	1
1	1	85	66	29	0	26.6	0.351	0
2	8	183	64	0	0	23.3	0.672	1
3	1	89	66	23	94	28.1	0.167	0
4	0	137	40	35	168	43.1	2.288	1

```

In [4]: healthcare.info()

```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 768 entries, 0 to 767
Data columns (total 9 columns):
#   Column                Non-Null Count  Dtype  
---  -
0   Pregnancies            768 non-null   int64  
1   Glucose                 768 non-null   int64  
2   BloodPressure          768 non-null   int64  
3   SkinThickness          768 non-null   int64  
4   Insulin                768 non-null   int64  
5   BMI                    768 non-null   float64 
6   DiabetesPedigreeFunction 768 non-null   float64 
7   Age                    768 non-null   int64  
8   Outcome                768 non-null   int64  
dtypes: float64(2), int64(7)
memory usage: 54.1 KB

```

Approach:

1. Perform descriptive analysis. It is very important to understand the variables and corresponding values. We need to think through - Can minimum value of below listed columns be zero (0)? On these columns, a value of zero does not make sense and thus indicates missing value.

- Glucose
- BloodPressure
- SkinThickness

- Insulin
- BMI

In [5]: `healthcare.describe()`

Out[5]:

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPe
count	768.000000	768.000000	768.000000	768.000000	768.000000	768.000000	
mean	3.845052	120.894531	69.105469	20.536458	79.799479	31.992578	
std	3.369578	31.972618	19.355807	15.952218	115.244002	7.884160	
min	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	
25%	1.000000	99.000000	62.000000	0.000000	0.000000	27.300000	
50%	3.000000	117.000000	72.000000	23.000000	30.500000	32.000000	
75%	6.000000	140.250000	80.000000	32.000000	127.250000	36.600000	
max	17.000000	199.000000	122.000000	99.000000	846.000000	67.100000	

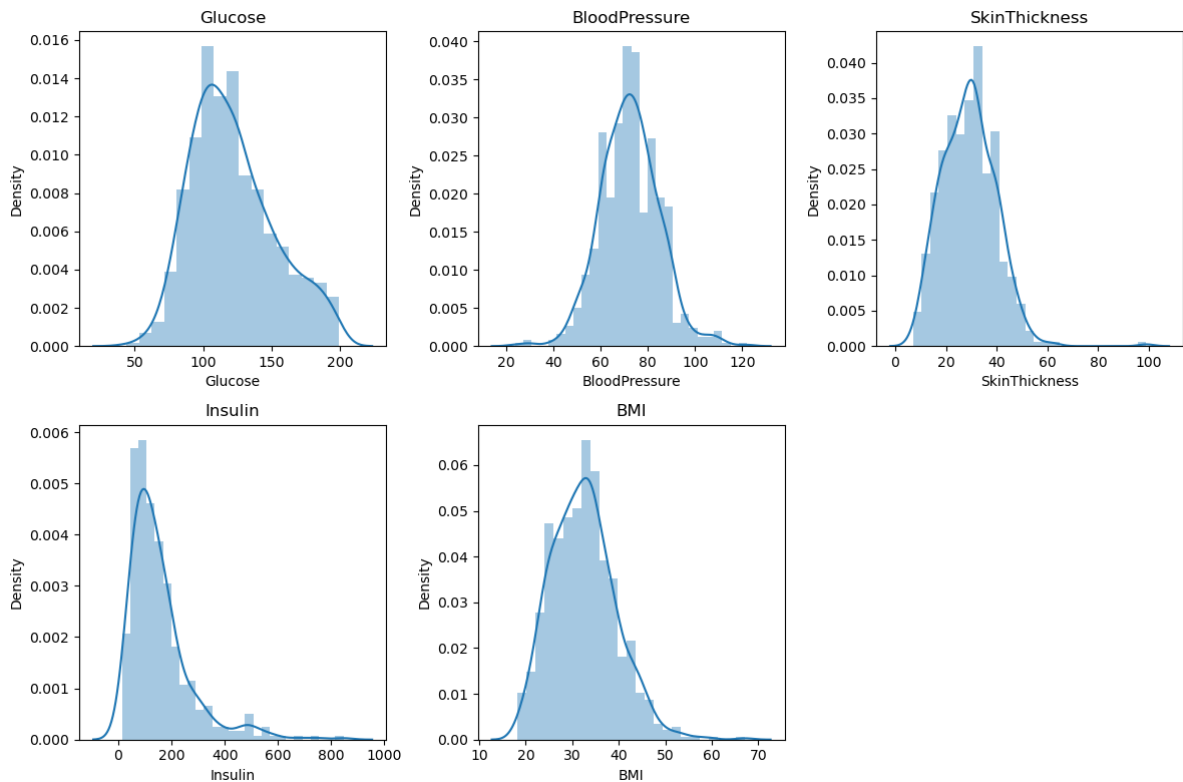
Descriptive analysis shows the minimum value as zero for variables Glucose, Bloodpressure, Skin thickness, Insulin and BMI. Where as these values cannot be zero, it indicates missing values and need to be treated. Hence transform these zero values to Nan values

In [6]: `# Columns List with zero values`
`col_list=['Glucose','BloodPressure','SkinThickness','Insulin','BMI']`

In [7]: `# Transforming the zero value to Nan values for the listed columns`
`for col in col_list:`
`healthcare.loc[healthcare[col]==0,[col]]=np.nan`

2. Visually explore these variable, you may need to look for the distribution of these variables using histograms. Treat the missing values accordingly.

In [8]: `plt.figure(figsize=(12,8))`
`for i,col in enumerate(col_list):`
`plt.subplot(2,3,i+1)`
`sns.distplot(healthcare.loc[:,col])`
`plt.title(col)`
`plt.tight_layout();`



From the above plot, it shows that the data is not normally distributed hence imputing the missing value (or zero values) with its corresponding median value

```
In [9]: # Imputing the Nan values with its corresponding median Values
for col in col_list:
    healthcare[col].fillna(healthcare[col].median(), inplace=True)
```

```
In [10]: ## Checking null values
healthcare.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 768 entries, 0 to 767
Data columns (total 9 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Pregnancies                          768 non-null    int64
1   Glucose                              768 non-null    float64
2   BloodPressure                        768 non-null    float64
3   SkinThickness                        768 non-null    float64
4   Insulin                              768 non-null    float64
5   BMI                                  768 non-null    float64
6   DiabetesPedigreeFunction             768 non-null    float64
7   Age                                  768 non-null    int64
8   Outcome                              768 non-null    int64
dtypes: float64(6), int64(3)
memory usage: 54.1 KB
```


3. We observe integer as well as float data-type of variables in this dataset. Create a count (frequency) plot describing the data types and the count of variables.

```
In [11]: healthcare.dtypes.value_counts()
```

```
Out[11]: float64    6
         int64     3
         dtype: int64
```

```
In [12]: sns.countplot(healthcare.count(),hue=healthcare.dtypes)
```

```
Out[12]: <AxesSubplot:ylabel='count'>
```



Above Plot indicates in the dataset 3 variables have Integer data type and 6 variables have Float data type

4. Check the balance of the data by plotting the count of outcomes by their value. Describe your findings and plan future course of actions.


```
In [13]: # Counting the outcomes by their value
healthcare.Outcome.value_counts()
```

```
Out[13]: 0    500
         1    268
         Name: Outcome, dtype: int64
```

```
In [14]: healthcare.Outcome.value_counts(normalize=True)
```

```
Out[14]: 0    0.651042
         1    0.348958
         Name: Outcome, dtype: float64
```

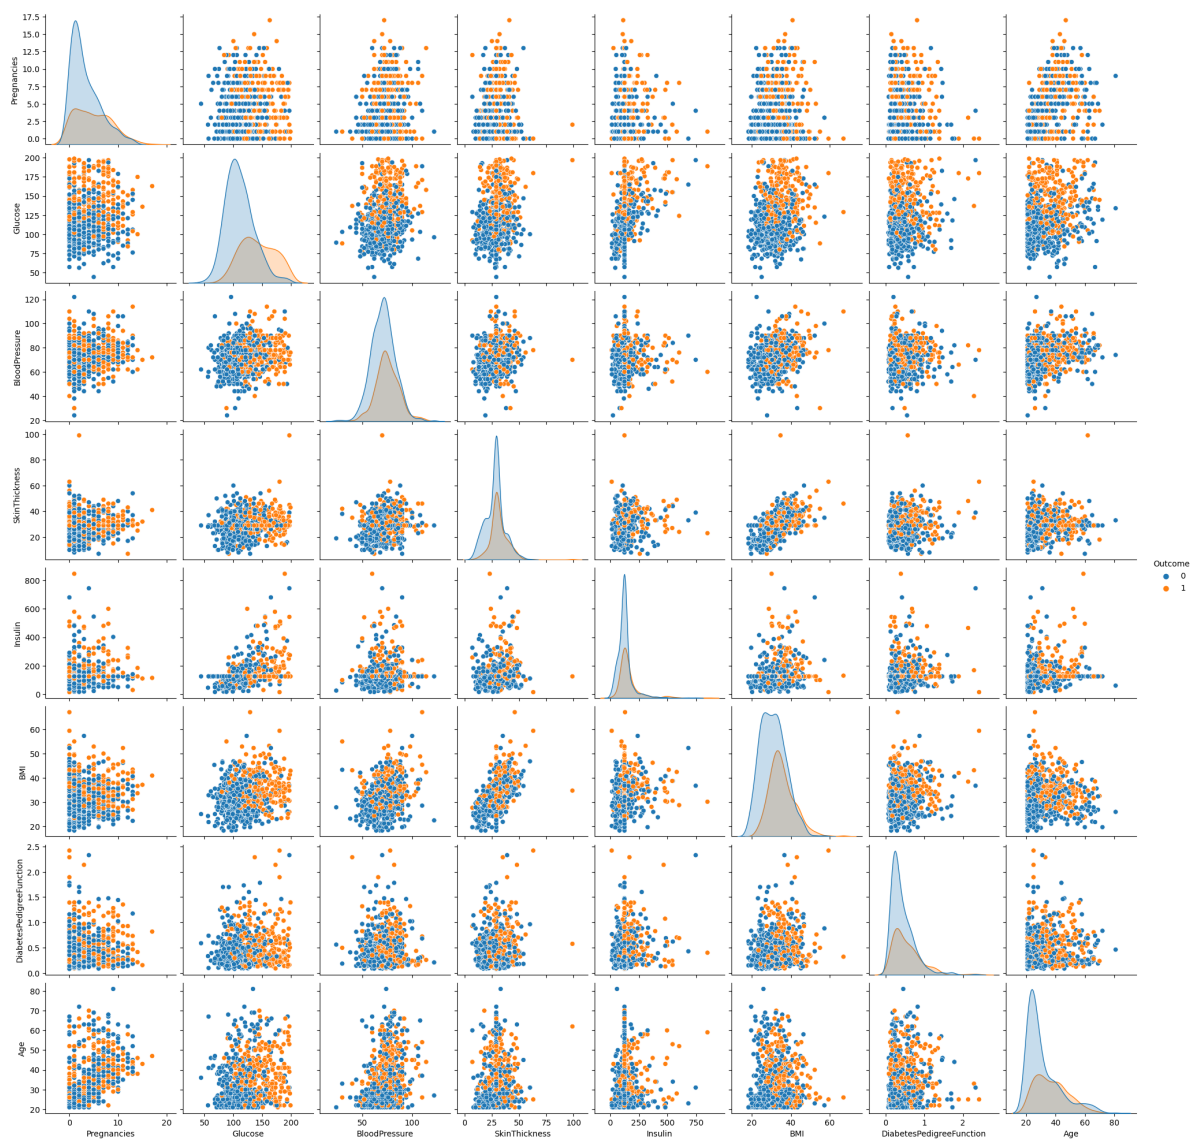
```
In [15]: sns.countplot(healthcare['Outcome'])
plt.title('Count of Outcome');
```



The above plot shows that the target variable (Outcome) is unevenly distributed and indicates data imbalanced, where 65% are non diabetic and 35% are diabetic. Need to treat data using SMOTE technique before building the model

5. Create scatter charts between the pair of variables to understand the relationships. Describe your findings.


```
In [16]: sns.pairplot(healthcare, kind='scatter', hue='Outcome')  
plt.show()
```



Above Scatter Plots indicates all independent (input) variables have weak linear correlation, and the diagonal plots show the data distribution.

6. Perform correlation analysis. Visually explore it using a heat map.

```
In [17]: sns.heatmap(healthcare.corr(), annot=True);
```



Above fig shows, all the independent variables correlation coefficient varies between -0.034 to 0.54 which displays weak correlation between variables

7. Devise strategies for model building. It is important to decide the right validation framework. Express your thought process. Would Cross validation be useful in this scenario?

1. Apply SMOTE technique to balance data and then prepare train and test data set
2. Apply ML model such as logistic regression, KNN model, Support Vector Machine, decision tree, random forest, Gradient boost and adaboost
3. Cross validation technique will help in building better model by lowering the bias and variances and using Grid search CV to tune the hyperparameters to obtain the optimum accuracy
4. Then compare the model accuracy score of each model. Choose the Model with the best accuracy score, Create confusion matrices, the Classification report, AUC(ROC curve)

```
In [18]: #Defining X and Y variables for the Model
X= healthcare.drop(['Outcome'],axis=1)
Y=healthcare['Outcome']
```

```
In [19]: X.shape , Y.shape
```


Out[19]: ((768, 8), (768,))

```
In [20]: # Scaling the X Variables
scaler=StandardScaler()
X_scaled=pd.DataFrame(scaler.fit_transform(X),columns=X.columns)
```

```
In [21]: X_scaled.describe().T
```

```
Out[21]:
```

	count	mean	std	min	25%	50%	75%
Pregnancies	768.0	2.544261e-17	1.000652	-1.141852	-0.844885	-0.250952	0.639947
Glucose	768.0	7.661695e-18	1.000652	-2.552931	-0.720163	-0.153073	0.611265
BloodPressure	768.0	-1.123956e-17	1.000652	-4.002619	-0.693761	-0.031990	0.629782
SkinThickness	768.0	-1.795800e-16	1.000652	-2.516429	-0.467597	-0.012301	0.329171
Insulin	768.0	4.416317e-17	1.000652	-1.467353	-0.222085	-0.181541	-0.155477
BMI	768.0	2.815312e-16	1.000652	-2.074783	-0.721209	-0.022590	0.603256
DiabetesPedigreeFunction	768.0	2.462585e-16	1.000652	-1.189553	-0.688969	-0.300128	0.466227
Age	768.0	1.857600e-16	1.000652	-1.041549	-0.786286	-0.360847	0.660206

Applying SMOTE to balance the Data

```
In [22]: from imblearn.over_sampling import SMOTE
```

```
In [23]: smote=SMOTE(sampling_strategy='minority')
X_sm,y_sm=smote.fit_resample(X_scaled,Y)
y_sm.value_counts()
```

```
Out[23]: 1    500
0    500
Name: Outcome, dtype: int64
```

Data is balanced considering the parameter sampling strategy as 'minority' to resample only the minority class that is diabetic

```
In [24]: # Train and Test data split
train_X,test_X,train_Y,test_Y=train_test_split(X_sm,y_sm,test_size=0.3,random_state=42)
```

```
In [25]: train_X.shape,test_X.shape,train_Y.shape,test_Y.shape
```

```
Out[25]: ((700, 8), (300, 8), (700,), (300,))
```

```
In [26]: train_Y.value_counts() , test_Y.value_counts()
```

```
Out[26]: (1    350
          0    350
          Name: Outcome, dtype: int64,
          0    150
          1    150
          Name: Outcome, dtype: int64)
```

a. Logistic Regression Model

```
In [27]: log_reg=LogisticRegression()
log_reg.fit(train_X,train_Y)
metrics.accuracy_score(test_Y,log_reg.predict(test_X))
```

```
Out[27]: 0.7266666666666667
```

```
In [28]: cross_val_score(log_reg,X_sm,y_sm,cv=5).mean()
```

```
Out[28]: 0.7550000000000001
```

```
In [29]: Model_Performances= dict()
```

```
In [30]: Model_Performances["logistic regression"]=cross_val_score(log_reg,X_sm,y_sm,cv=5).mean()
```

b. KNN Classifier


```
In [31]: knn=KNeighborsClassifier()
knn.fit(train_X,train_Y)
print('Test Score:{}'.format(knn.score(test_X,test_Y)))
```

```
Test Score:0.7733333333333333
```

```
In [32]: acc=list()
for n in range(2, int(np.sqrt(train_X.shape[0]))):
    knn= KNeighborsClassifier(n_neighbors=n)
    knn.fit(train_X,train_Y)
    acc.append(cross_val_score(knn,X_sm,y_sm,cv=5).mean())
```

```
In [33]: plt.plot(range(2, int(np.sqrt(train_X.shape[0]))),acc)
```

```
Out[33]: [<matplotlib.lines.Line2D at 0x1fb1ac45490>]
```



```
In [36]: knn= KNeighborsClassifier(n_neighbors=4)
knn.fit(train_X,train_Y)
knn.score(test_X,test_Y)
```

Out[36]: 0.77

```
In [37]: cross_val_score(knn,X_sm,y_sm,cv=5).mean()
```

Out[37]: 0.7919999999999999

```
In [38]: Model_Performances['K Nearest Neighbour']=cross_val_score(knn,X_sm,y_sm,cv=5).mean
```

c. Support Vector Machine

```
In [39]: # Tuning hyperparameter for SVM using grid search
param_grid={'C':[0.001,0.01,0.1,1,10], 'gamma':[0.001,0.01,0.1,1,10]}
gscv=GridSearchCV(SVC(),param_grid,cv=5,verbose=1)
gscv.fit(X_sm,y_sm)
```

Fitting 5 folds for each of 25 candidates, totalling 125 fits

```
Out[39]: GridSearchCV(cv=5, estimator=SVC(),
    param_grid={'C': [0.001, 0.01, 0.1, 1, 10],
    'gamma': [0.001, 0.01, 0.1, 1, 10]},
    verbose=1)
```

```
In [40]: gscv.best_estimator_
```

Out[40]: SVC(C=10, gamma=1)

```
In [41]: cross_val_score(gscv.best_estimator_,X_sm,y_sm,cv=5).mean()
```

Out[41]: 0.8399999999999999

```
In [42]: Model_Performances['Support Vector Machine']=cross_val_score(gscv.best_estimator_,
```

d. Decision Tree Classifier

```
In [43]: param_grid={'max_depth':[3,5,7,None], 'min_samples_leaf':[3,5,7,9]}
gscv_DT=GridSearchCV(DecisionTreeClassifier(),param_grid,cv=5,verbose=1)
gscv_DT.fit(X_sm,y_sm)
```

Fitting 5 folds for each of 16 candidates, totalling 80 fits

```
Out[43]: GridSearchCV(cv=5, estimator=DecisionTreeClassifier(),
      param_grid={'max_depth': [3, 5, 7, None],
      'min_samples_leaf': [3, 5, 7, 9]},
      verbose=1)
```

```
In [44]: gscv_DT.best_estimator_
```

```
Out[44]: DecisionTreeClassifier(min_samples_leaf=7)
```

```
In [45]: cross_val_score(gscv_DT.best_estimator_,X_sm,y_sm,cv=5).mean()
```

```
Out[45]: 0.764
```

```
In [46]: Model_Performances['Decision Tree Classifier']=cross_val_score(gscv_DT.best_estimator_,X_sm,y_sm,cv=5).mean()
```

e. Random Forest Classifier

```
In [47]: param_grid={ 'n_estimators':[50,100,150,200], 'max_depth':[3,5,7,None], 'min_samples_leaf':[3,5,7,9]}
gscv_RFC=GridSearchCV(RandomForestClassifier(),param_grid,cv=5,verbose=1)
gscv_RFC.fit(X_sm,y_sm)
```

Fitting 5 folds for each of 64 candidates, totalling 320 fits

```
Out[47]: GridSearchCV(cv=5, estimator=RandomForestClassifier(),
      param_grid={'max_depth': [3, 5, 7, None],
      'min_samples_leaf': [3, 5, 7, 9],
      'n_estimators': [50, 100, 150, 200]},
      verbose=1)
```

```
In [48]: gscv_RFC.best_estimator_
```

```
Out[48]: RandomForestClassifier(min_samples_leaf=3, n_estimators=150)
```

```
In [49]: cross_val_score(gscv_RFC.best_estimator_,X_sm,y_sm,cv=5).mean()
```

```
Out[49]: 0.8039999999999999
```

```
In [50]: Model_Performances['Random Forest Classifier']=cross_val_score(gscv_RFC.best_estimator_,X_sm,y_sm,cv=5).mean()
```

f. Gradient Boosting

```
In [51]: GBC=GradientBoostingClassifier()
GBC.fit(train_X,train_Y)
GBC.score(test_X,test_Y)
```

```
Out[51]: 0.7833333333333333
```

```
In [52]: cross_val_score(GBC,X_sm,y_sm,cv=5).mean()
```

```
Out[52]: 0.7929999999999999
```

```
In [53]: Model_Performances['Gradient Boosting Classifier']=cross_val_score(GBC,X_sm,y_sm,cv
```

g. Adaptive Boosting classifier

```
In [54]: param_grid={ 'n_estimators':[50,100,150,200],'base_estimator':[DecisionTreeClassif:
gscv_AB=GridSearchCV(AdaBoostClassifier(),param_grid,cv=5,verbose=1)
gscv_AB.fit(X_sm,y_sm)
```

Fitting 5 folds for each of 8 candidates, totalling 40 fits

```
Out[54]: GridSearchCV(cv=5, estimator=AdaBoostClassifier(),
param_grid={'base_estimator': [DecisionTreeClassifier(max_depth=5),
DecisionTreeClassifier(max_depth=5)],
'n_estimators': [50, 100, 150, 200]},
verbose=1)
```

```
In [55]: gscv_AB.best_estimator_
```

```
Out[55]: AdaBoostClassifier(base_estimator=DecisionTreeClassifier(max_depth=5),
n_estimators=200)
```

```
In [56]: cross_val_score(gscv_AB.best_estimator_,X_sm,y_sm,cv=5).mean()
```

```
Out[56]: 0.8300000000000001
```

```
In [57]: Model_Performances['Adaptive Boosting Classifier']=cross_val_score(gscv_AB.best_es
```

8. Apply an appropriate classification algorithm to build a model. Compare various models with the results from KNN.


```
In [58]: # Compare all the Model Performances
Model_performance=pd.DataFrame(Model_Performances.items(),columns=['Classifier-Model',
Model_performance
```

```
Out[58]:
```

	Classifier-Model	Model_Performance
0	logistic regression	0.755
1	K Nearest Neighbour	0.792
2	Support Vector Machine	0.840
3	Decision Tree Classifier	0.760
4	Random Forest Classifier	0.810
5	Gradient Boosting Classifier	0.794
6	Adaptive Boosting Classifier	0.821

	Classifier-Model	Model_Performance
0	logistic regression	0.755
1	K Nearest Neighbour	0.792
2	Support Vector Machine	0.840
3	Decision Tree Classifier	0.760
4	Random Forest Classifier	0.810
5	Gradient Boosting Classifier	0.794
6	Adaptive Boosting Classifier	0.821

```
In [59]: sns.barplot(data=Model_performance,x='Classifier-Model',y='Model_Performance')
plt.xticks(rotation=90)
plt.grid(True)
```



Comparison of the model performances indicates, Support Vector Machine, Random forest and Adaboost accuracy score is above 0.8 with a highest score of 0.84 for support vector machine


Final Classifier model- Support Vector Machine

```
In [60]: svc=SVC(C=10, gamma=1)
svc.fit(train_X,train_Y)
svc.score(test_X,test_Y)
```

```
Out[60]: 0.8366666666666667
```

9. Create a classification report by analysing sensitivity, specificity, AUC(ROC curve) etc. Please try to be as descriptive as possible to explain what values of these parameter you settled for? any why?

```
In [61]: # Confusion Matrix
sns.heatmap(metrics.confusion_matrix(test_Y,svc.predict(test_X)),annot=True);
```



```
In [62]: print(metrics.classification_report(test_Y,svc.predict(test_X)))
```

	precision	recall	f1-score	support
0	0.86	0.80	0.83	150
1	0.81	0.87	0.84	150
accuracy			0.84	300
macro avg	0.84	0.84	0.84	300
weighted avg	0.84	0.84	0.84	300

The Accuracy score of the model is 84% and with f1 score of 0.84

```
In [63]: y_pred = svc.decision_function(test_X)
fpr_svc, tpr_svc, _ = roc_curve(test_Y, y_pred)
roc_auc_svc = auc(fpr_svc, tpr_svc)
print("Area under ROC curve = {:.2f}".format(roc_auc_svc))
```

Area under ROC curve = 0.91

AUC value measures the degree of separability in the classification model. In the above calculation AUC value is 0.91 which is between 0.9-1 which is considered as good score

```
In [64]: # Transferring the healthcare dataframe to excel for preparing the report in Tableau
healthcare.to_excel('healthcare.xlsx')
```